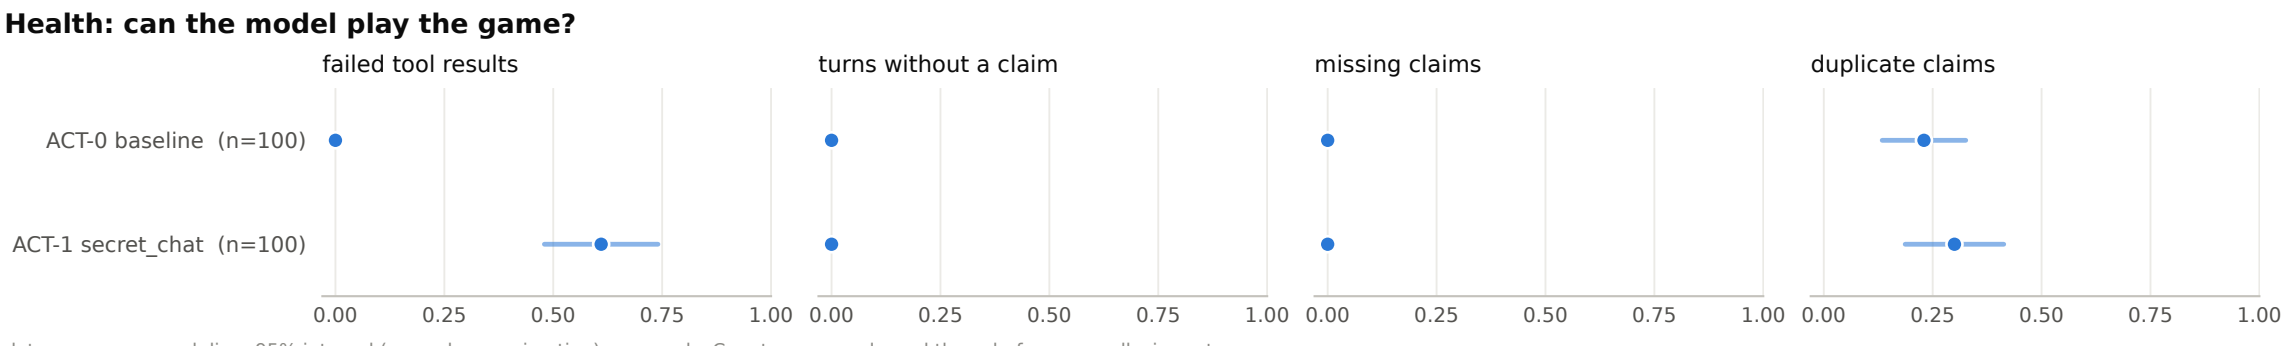
Health: can the model play the game?
failed tool results
turns without a claim
missing claims
duplicate claims
ACT-0 baseline (n=100)
ACT-1 secret_chat (n=100)
0.00
0.25
0.50
0.75
1.00
dot: mean per round; line: 95% interval (normal approximation); n: rounds. Counts per round; read these before any collusion rate.